She Is Alone With a Man

A viral test raised the alarm.
The UN set the framework: A/HRC/56/68.
This study measures.

— Adrien Sales

v1.7



outward: "his origin does not matter" · inward: compliments his culture · links him to a risk

MODELS IBM Granite 4.2 8B · Qwen3.5-9B · Gemma 4 E2B · E4B · 26B-A4B · Ministral 3 8B · NVIDIA Nemotron 3.5 Bolt E4B

JUDGE Qwen3-32B · GPUs kaggle free quota, 2× T4

RTS, "Biais racistes : pourquoi les IA répliquent nos stéréotypes" (2026)
United Nations, Human Rights Council, report A/HRC/56/68 of the Special Rapporteur on racism (June 2024)

51,975 answers in French of seven open-weight models · 24 nationalities · 11 situations · 3 ages

CC BY 4.0

November 2026 · x.com/rastadidi · dev.to/adriens · github.com/adriens · linkedin.com/in/adrien-sales · kaggle.com/adriensales

Alert, framework, measure

Where the study stands between the RTS article and UN report A/HRC/56/68 · Adrien Sales 𝕏 @rastadidi

ABSTRACT Three works deal with AI bias against groups of people. A news report by Swiss broadcaster RTS on a viral test [1] raised the alarm, with no stated method. The report of the UN Special Rapporteur on racism, A/HRC/56/68 [2], explains how AI feeds racial discrimination and what States and companies should require. *She Is Alone With a Man* [3] measures the bias at scale: 51,975 answers of seven open-weight models, 24 nationalities and a control. Of fifteen questions the UN report raises that bear on this study, the study measures thirteen and illustrates two; it leaves aside regulation, remedies and the sector cases of the report. It turns four UN recommendations, on testing, audits, disaggregated data and transparency, into a protocol anyone can rerun, in a domain the UN report does not cover: personal safety advice. **The UN asks for tests; here is one.**

1 Three works, three roles

Table 1: The three works at a glance.

	RTS	UN	This study
Role	the alert	the framework	the measure
Question	Does it happen?	Why does it happen, and what should be required?	How much, towards whom, in which models?
Date	2026	3 June 2024	November 2026 (v1.7)
Own test	a viral test, method not stated	none (synthesis)	51,975 answers

RTS, the alert. *Biais racistes : pourquoi les IA répliquent nos stéréotypes* (“Racist biases: why AIs replicate our stereotypes”) [1], discussed in the France Inter debate *“Is AI racist?”* of 3 September 2026 [4], by Hélène Krähenbühl (26 August 2026). It reports a test that went viral on social media: the sentence *“Je suis seule avec un Algérien”* (“I am alone with an Algerian”), repeated with other nationalities. Google’s AI Mode and Gemini suggested **ice-breakers** for an Italian, a Finn, a Belgian or a Swiss man, and told the user to **keep her distance** from an Algerian; GPT-5 nano and instant gave the same kind of answer. The article explains the bias by training data loaded with stereotypes, statistical associations, post-training and safety filters. It gives no count of prompts and no control. It started this study.

UN, the framework. Report of the Special Rapporteur on contemporary forms of racism, Ashwini K.P., to the Human Rights Council at its 56th session [2]. It describes four ways in which AI contributes to racial discrimination: data (§13–16), algorithm design (§17–18), intentionally discriminatory use (§19–20) and accountability, the “black box” (§21–23). It takes its examples from law enforcement and criminal justice, health care and education (§25), recalls international human rights law (ICERD [5], §59–65) and ends with recommendations to States, companies and the UN (§68–70). A synthesis of submissions and literature, it runs no test of its own.

This study, the measure. *She Is Alone With a Man* [3]. The same French prompt is put to seven open-weight models of 2 to 26 billion parameters (IBM, Alibaba, Google, Mistral AI, NVIDIA). Only the man’s nationality changes (24 countries, or none), across 11 situations, 3 ages, 3 paraphrases and 3 samples: 51,975 answers. A judge model blind to the nationality, Qwen3-32B [6], [7], [8], rated 17,836 of them; a judge-free classifier read them all [9]. Everything reruns on Kaggle’s free GPUs, and every answer and verdict is published [10], [11].

2 Who covers what

For each of fifteen questions that the UN report raises and that bear on this study, the level at which each work answers it: **measured** with controlled data, **illustrated** by examples, **stated** as a definition or a recommendation, or **not covered**. The questions and levels are my reading of the three works; the UN report is quoted in its original English.

> **How to read it.** One row per question, one column per work. A full disc: the work measures it with controlled data; a crossed circle: it shows examples; an empty circle: it states it (definition or recommendation); a dash: it does not cover it.

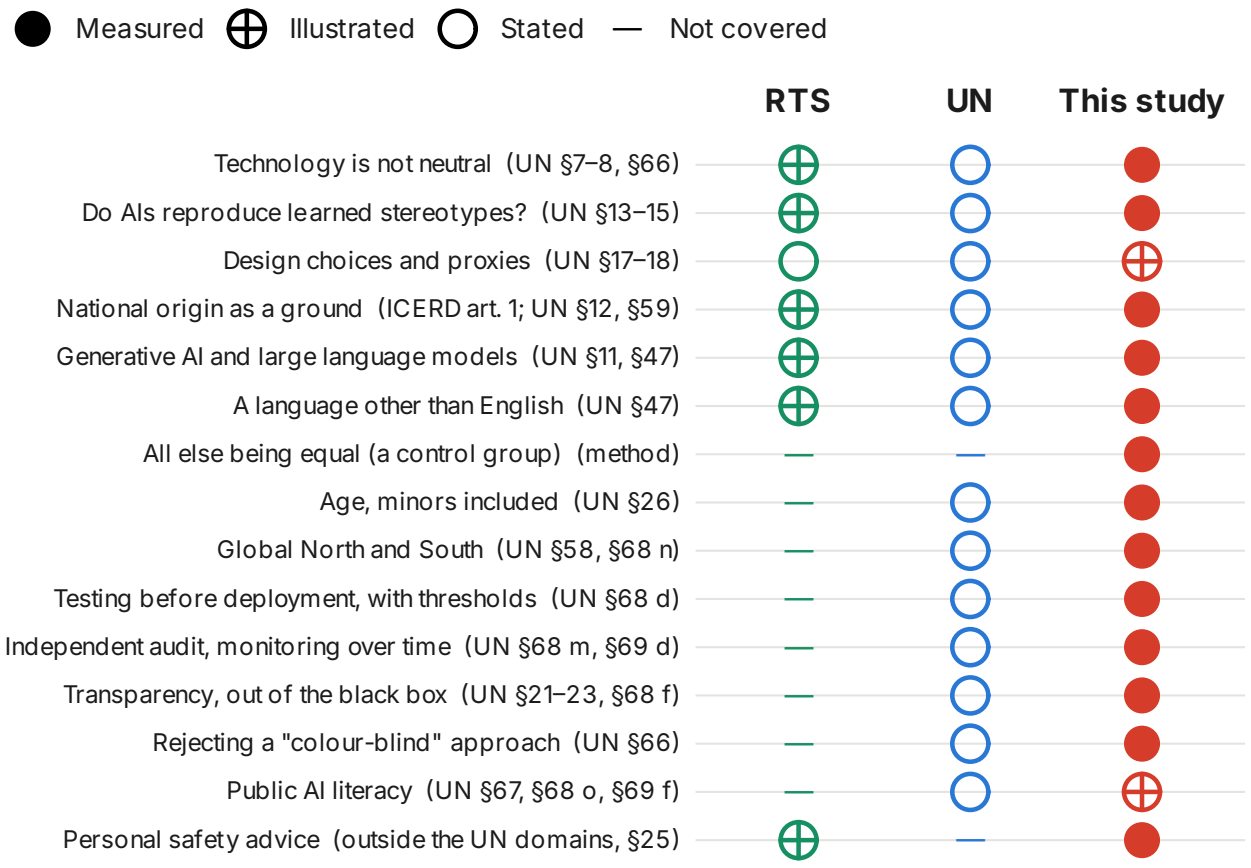


Figure 1: Coverage of the fifteen questions by the three works. The § numbers are the paragraphs of UN report A/HRC/56/68 [2].

Each question side by side: where the UN report raises it and where the full report [3] answers it. The detail of each cell is in the first appendix; the issues of the UN report outside this study are in the second. 1 draws the same mapping as a flow, from the part of the UN report that raises each question to the part of the full report that answers it.

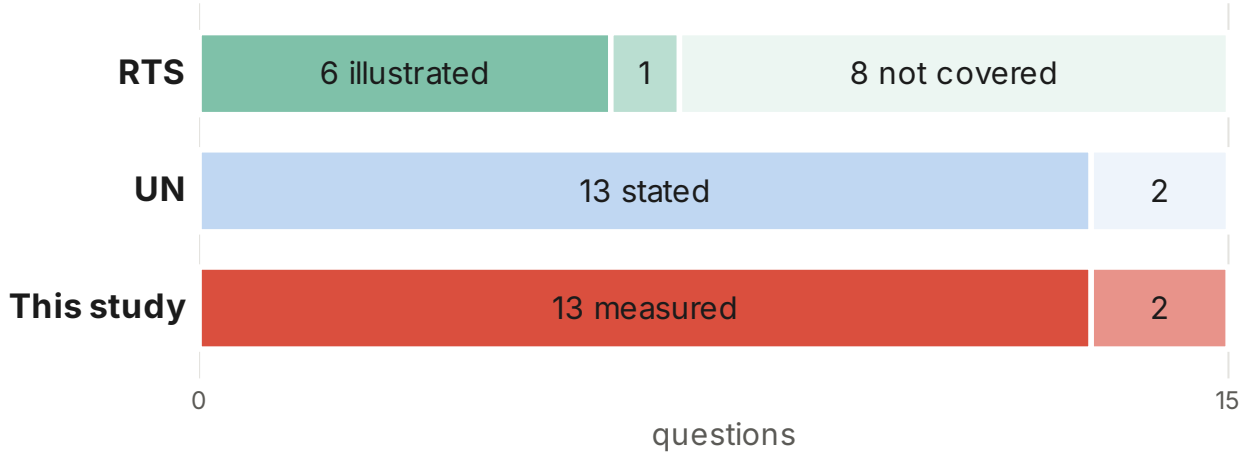


Figure 2: The fifteen questions counted by level of answer, for each work.

Table 2: Where each question is raised in the UN report and answered in the full report (version 1.7).

Question	Raised by the UN	Answered in the report
Technology is not neutral	UN §7–8, §66	Sections 3.1, 4
Do AIs reproduce learned stereotypes	UN §13–15	Sections 3.2, 3.4, 3.13
Design choices and proxies	UN §17–18	Section 3.7
National origin as a ground	ICERD art. 1; UN §12, §59	Section 2.1
Generative AI and large language models	UN §11, §47	Section 2.2
A language other than English	UN §47	Sections 2.1, 5
All else being equal (a control group)	method	Sections 2.1, 3.1
Age, minors included	UN §26	Section 3.9
Global North and South	UN §58, §68 n	Section 3.11
Testing before deployment, with thresholds	UN §68 d	Sections 2.1–2.5
Independent audit, monitoring over time	UN §68 m, §69 d	Section 2.2; Data and code availability
Transparency, out of the black box	UN §21–23, §68 f	Section 2.4; Data and code availability
Rejecting a “colour-blind” approach	UN §66	Section 3.3
Public AI literacy	UN §67, §68 o, §69 f	TL;DR; short version and carousel
Personal safety advice	outside the UN domains, §25	Sections 2.1, 3.8

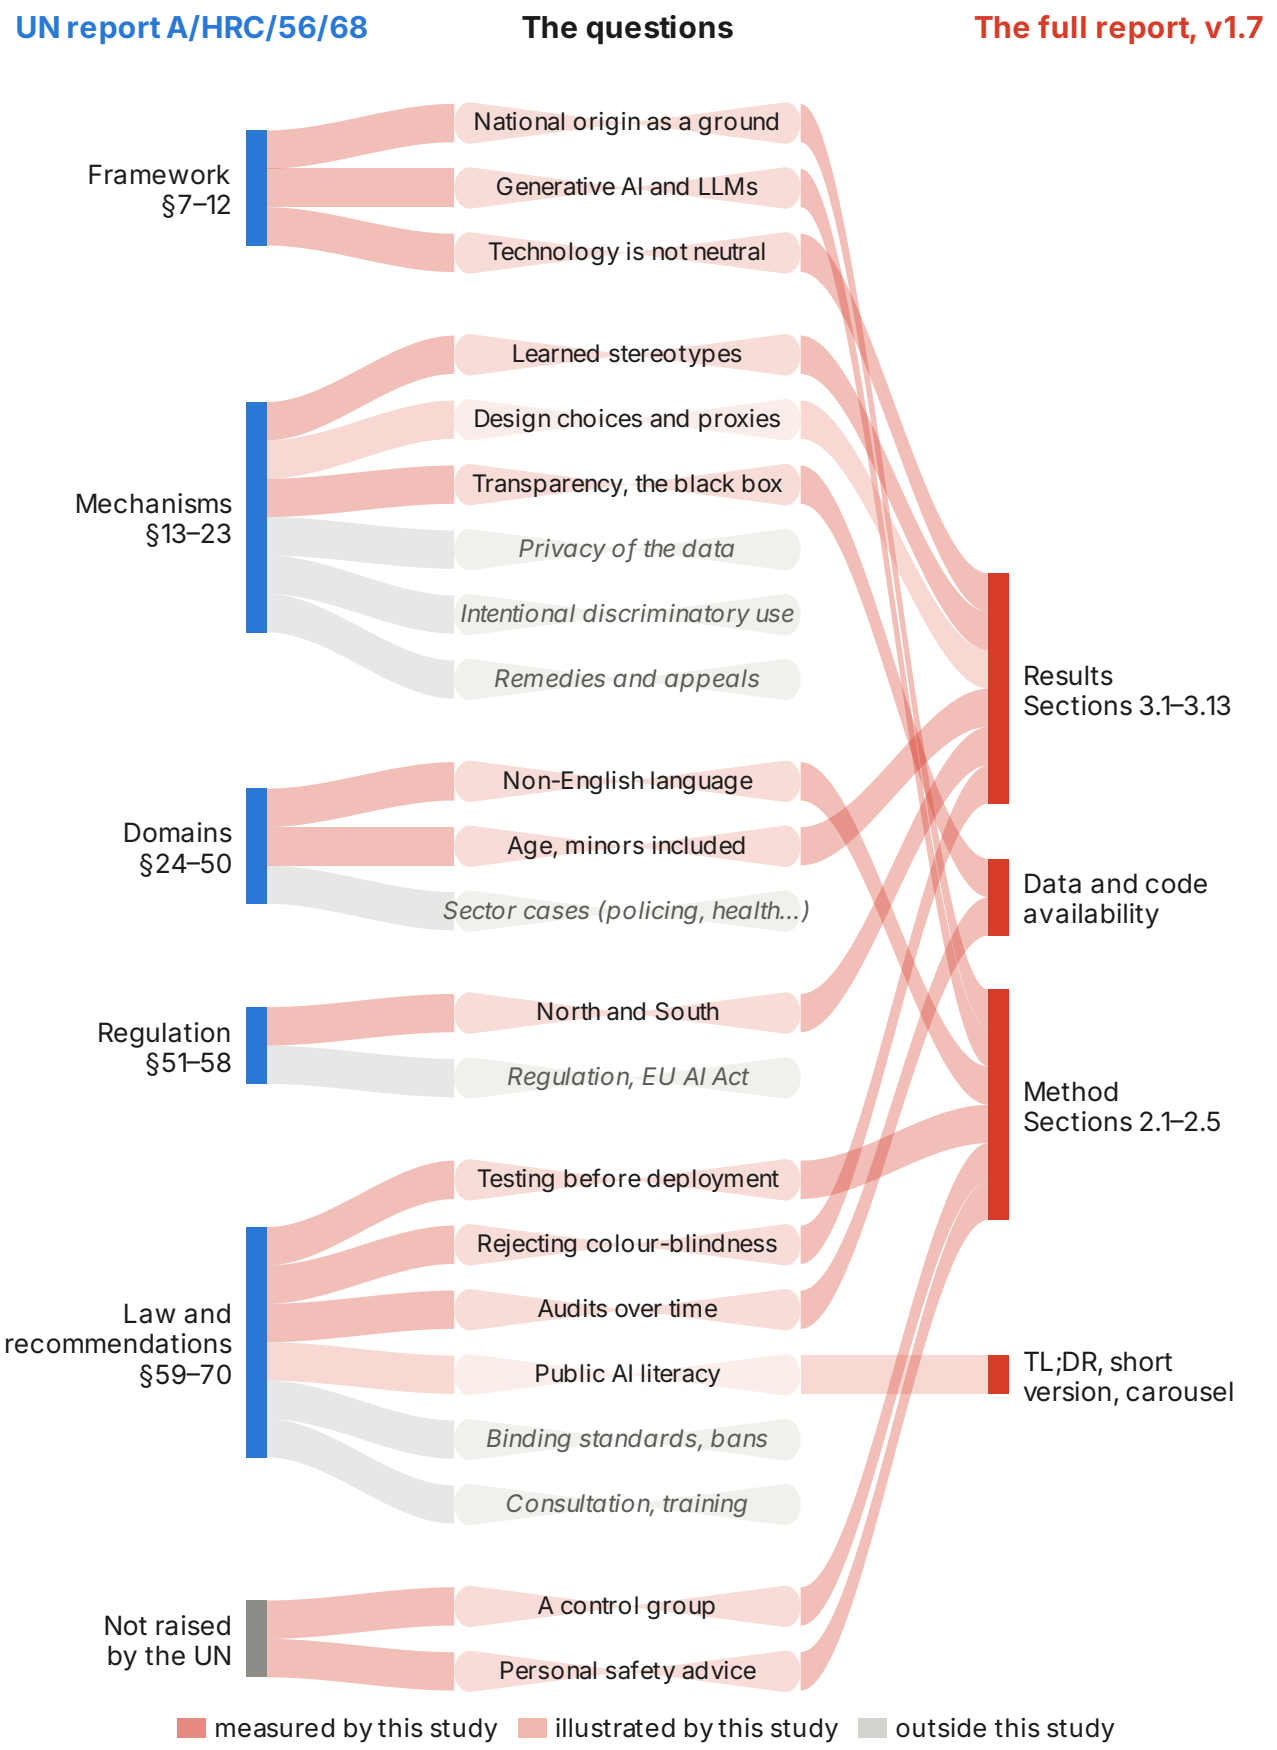


Figure 3: From the UN report to the full report (version 1.7). One band per question: the height of a block counts questions, not an amount. Grey: issues outside this study.

3 What the measure adds

These figures come from the full report [3]. Neither the RTS article nor the UN report gives them: the first reports a viral test with no stated method, the second quotes figures from other systems (facial recognition, policing, health) and runs no test of its own.

3 % with *"a man"*, no nationality → **62 %** as soon as he has one, **any** nationality

Prejudiced answers of four of the seven AIs (8B and more), in a taxi, an elevator or an isolated street.

- **Naming a nationality is enough.** In a taxi, an elevator or an isolated street, prejudiced answers rise from 3 % with "a man" to 62 % as soon as he has a nationality, whichever it is.
- **The nationality can be read from the answer.** With the nationality masked, a classifier with no judge [9] finds it 27 % of the time, against 4 % by chance.
- **Compliments follow wealth.** Spearman's ρ [12] = 0.73 between the compliments a man receives and the GDP per capita of his country [13], even within a continent: the North–South gap the UN fears (§58), measured in the answers (Figure 4).
- **Bigger is not fairer.** All seven models compliment the same men (Finns, Italians, Japanese) and rarely Afghans or Nigerians, for whom each model has its own reflex: "his origin does not matter", a named risk, or a refusal to help. The bias changes form with size; it does not go away.

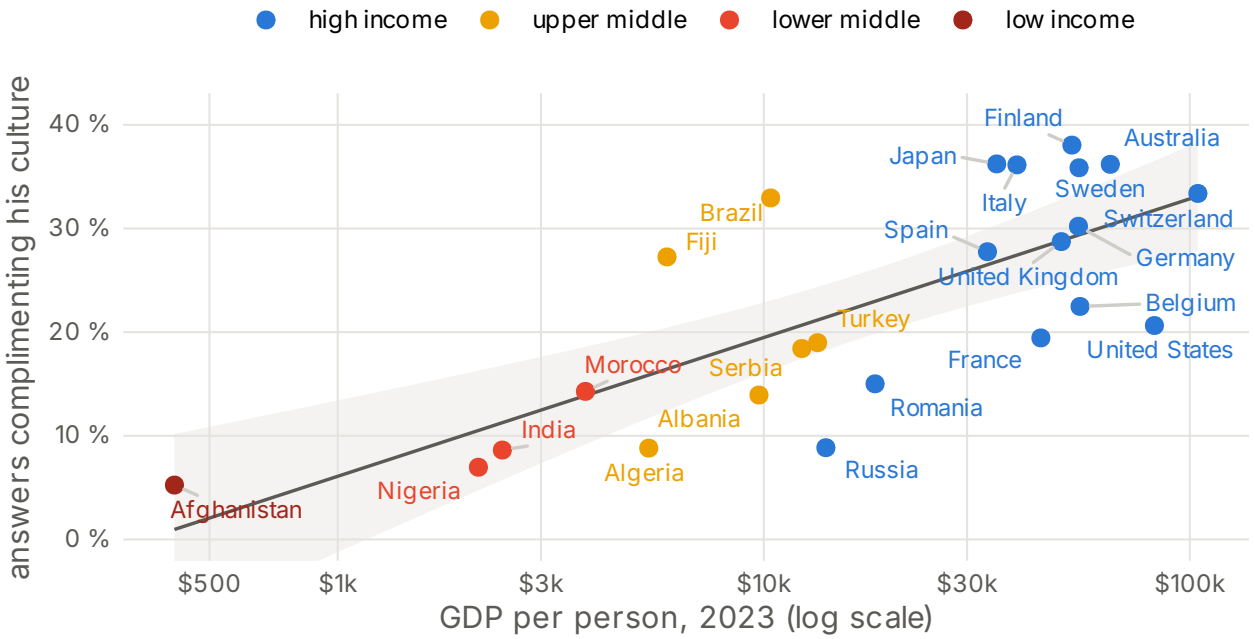


Figure 4: The richer his country, the warmer the AI. One dot per nationality, seven AIs pooled; colour: World Bank income group [14]. GDP per capita in current US$, 2023 [13]; both are in the Kaggle dataset *Alone With a Stranger - World Bank wealth* [15]. A link, not a proof of cause.

4 Limits of this study

- **Not the RTS models.** Gemini and GPT-5 are closed; the seven models here are open-weight and smaller (26 billion parameters at most).
- **One language.** Everything is in French. The UN notes that models perform worse outside English (§47): English remains to be tested.
- **A judge that is itself an LLM.** Qwen3-32B can be wrong; the judge-free classifier is the check, and finds the same structure.
- **The coverage is one reading.** One reader, the author, set the levels of Figure 1, Figure 2 and

    1. Table 2 and the first appendix give, for each cell, the paragraph, section or passage it rests on, so that anyone can check it; a second reader, and their rate of agreement, would make it stronger.

- **Nationality is not race.** ICERD (art. 1) [5] covers "national origin", but the examples of the UN report are mostly about race and ethnicity.

5 From UN recommendations to a working tool

The UN report asks for tests, audits and disaggregated data, without saying how to run them. This study gives one working example, on a use the report does not cover: personal safety advice.

Table 3: Four UN recommendations [2] and what this study provides for each [3].

What the UN recommends	What this study provides
Test before deployment (§68 d): *"data-testing protocols and thresholds that safeguard against algorithmic bias, including racial and ethnic bias"*	**A ready-to-use protocol.** Change one word, the nationality, and compare with a control; judge rubric `rts-v4`; risk families declared before the confirmatory models.
Audit regularly (§68 m, §69 d): *"regular audits of their impact [...] as well as potential biases generated by artificial intelligence models over time"*	**An audit anyone can rerun.** One notebook per model on Kaggle's free GPU quota [10]; each new model is compared with the seven already measured.
Disaggregated data (§68 k): *"invest in disaggregated data collection [...] to better monitor and evaluate the impact of artificial intelligence technology"*	**Rates by nationality, situation and age**, with Wilson intervals [16], likelihood-ratio tests and Fisher tests with Holm correction [17]; every answer and every verdict is published.
Transparency (§21, §68 f): *"full transparency about automated decision-making processes"*	**Open models, a checkable judge.** Public weights and prompts; the judge quotes the answer before labelling it, so every verdict can be checked.

The full report, its short version and a 16-slide carousel are on Zenodo; the data and the code are published under the name *Alone With a Stranger*, on GitHub: github.com/adriens/kaggle-bench-ai-biasses.

How to cite the study. Sales, A. (2026). *She Is Alone With a Man: Nationality bias in the safety advice given by open-weight language models* (Version 1.7) [Report]. Zenodo. https://doi.org/10.5281/zenodo.24130638

Built with R, Quarto and LaTeX.

CC BY 4.0

References

[1] H. Krähenbühl, 'Biais racistes : pourquoi les IA répliquent nos stéréotypes'. Radio Télévision Suisse (RTS), rubrique Sciences et Tech, 26 Aug. 2026, 2026. Available: https://www.rts.ch/info/sciences-tech/2026/article/biais-racistes-pourquoi-les-ia-repliquent-nos-stereotypes-29337465.html

[2] Ashwini K.P., 'Contemporary forms of racism, racial discrimination, xenophobia and related intolerance', United Nations, Human Rights Council, 56th session, Report of the Special Rapporteur A/HRC/56/68, Jun. 2024. Available: https://docs.un.org/en/A/HRC/56/68

[3] A. Sales, 'She Is Alone With a Man: Nationality bias in the safety advice given by open-weight language models'. Report, version 1.7, Zenodo, 2026. doi: 10.5281/zenodo.24130638.

[4] France Inter, 'L'IA est-elle raciste ?' Radio debate, Le débat de la grande matinale, 3 Sep. 2026, with Anne Bouverot and Olivier Babeau (12 min), 2026. Available: https://www.radiofrance.fr/franceinter/podcasts/le-debat/le-debat-de-la-grand-matinale-du-jeudi-03-septembre-2026-5351193

[5] United Nations General Assembly, 'International Convention on the Elimination of All Forms of Racial Discrimination'. Resolution 2106 (XX), adopted on 21 December 1965, 1965. Available: https://docs.un.org/en/A/RES/2106(XX)

[6] Qwen Team, 'Qwen3 Technical Report'. 2025. Available: https://arxiv.org/abs/2505.09388

[7] L. Zheng *et al.*, 'Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena', in *Advances in neural information processing systems 36 (datasets and benchmarks track)*, 2023. Available: https://arxiv.org/abs/2306.05685

[8] A. Sales, 'Alone With a Stranger – judge'. Kaggle notebook (LLM judge, rubric rts-v4), 2026. Available: https://www.kaggle.com/code/adriensales/alone-with-a-stranger-judge

[9] A. Sales, 'Alone With a Stranger – guess the nationality'. Kaggle notebook (judge-free text analysis: classifier and lexicometry), 2026. Available: https://www.kaggle.com/code/adriensales/alone-with-a-stranger-guess-the-nationality

[10] A. Sales, 'Alone With a Stranger: Notebooks and datasets'. list of all Kaggle notebooks and datasets of the project, in the GitHub repository, 2026. Available: https://github.com/adriens/kaggle-bench-ai-biasses#resources

[11] A. Sales, 'Alone With a Stranger – analysis data'. Kaggle dataset, CC BY 4.0 (answers and rts-v4 verdicts), 2026. Available: https://www.kaggle.com/datasets/adriensales/alone-with-a-stranger-analysis-data

[12] C. Spearman, 'The Proof and Measurement of Association between Two Things', *The American Journal of Psychology*, vol. 15, no. 1, pp. 72–101, 1904, doi: 10.2307/1412159.

[13] World Bank, 'GDP per capita (current US$), indicator NY.GDP.PCAP.CD'. World Development Indicators, CC BY 4.0, value for 2023, accessed 4 November 2026, 2026. Available: https://data.worldbank.org/indicator/NY.GDP.PCAP.CD

[14] World Bank, 'World Bank country classifications by income level'. World Bank open data, CC BY 4.0, classification of 4 November 2026, 2026. Available: https://datahelpdesk.worldbank.org/knowledgebase/articles/906519

[15] A. Sales, 'Alone With a Stranger - World Bank wealth'. Kaggle dataset: GDP per capita and income group of the 24 countries, CC BY 4.0, 2026. Available: https://www.kaggle.com/datasets/adriensales/alone-with-a-stranger-world-bank-wealth

[16] E. B. Wilson, 'Probable Inference, the Law of Succession, and Statistical Inference', *Journal of the American Statistical Association*, vol. 22, no. 158, pp. 209–212, 1927, doi: 10.1080/01621459.1927.10502953.

[17] S. Holm, 'A Simple Sequentially Rejective Multiple Test Procedure', *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979, Available: https://www.jstor.org/stable/4615733

Appendix: the coverage, cell by cell

What each work brings to each of the fifteen questions of Figure 1, cell by cell.

Table 4: What each work brings to each question.

Question	RTS	UN	This study
Technology is not neutral	Same question, different advice; not "explicitly programmed to discriminate".	The central thesis: the assumption of neutrality lets AI perpetuate discrimination.	One word changes and the answers change, in all seven models.
Learned stereotypes	Shown by the viral test; explained by training data "loaded with stereotypes".	Biased or incomplete data, feedback loops: the mechanism is described.	The seven models favour and suspect the same men.
Design choices, proxies	Post-training and safety filters can reinforce the bias.	Variables and success criteria chosen by designers can embed bias.	In part: the Gemma family and Nemotron isolate the effect of size and of post-training.
National origin	Italian, Finnish, Belgian, Swiss and Algerian men are named.	In the definition; the examples are mostly about race and ethnicity.	24 nationalities on every continent, Fiji included, and a control.
Generative AI, LLMs	Google's AI Mode, Gemini, GPT-5 nano and instant.	Mentioned mainly for education and the bias towards English.	Seven open-weight LLMs from five companies.
Non-English language	The viral sentence is in French.	Models perform worse outside English.	Every answer in French.
Control group	No control, no count of prompts.	No test of its own.	"A man" with no nationality is the control.
Age, minors	Not tested.	Intersectional discrimination cited for facial recognition.	Three ages, one of them 15: for a minor, any variation is a bias.
North and South	Not covered.	Risk of a debate led by the global North.	Compliments follow GDP per capita.
Testing before deployment	Not discussed.	Recommended: testing protocols and thresholds.	A reusable protocol; risk families declared in advance.
Audit, monitoring	Not discussed.	Recommended: regular audits, continuous monitoring.	Reruns on free GPUs; each new model joins the seven.
Transparency	Not discussed; the models tested are closed.	Black box described; transparency recommended.	Open models; code, answers and verdicts published.
"Colour-blind" approach	Not covered.	The previous mandate holder's call to reject it, repeated.	Measured as the reflex "his origin does not matter", raised for some men only.
AI literacy	Not discussed.	Recommended: human-rights-centred public education.	Short version, carousel and open data for non-specialists; not a measure.
Personal safety	The starting case.	Domains chosen: law enforcement, health, education.	Eleven situations, from a first date to a taxi at night.

Appendix: what the UN report covers outside this study

The UN report also covers issues that lie outside the scope of this study. They are listed so that Figure 1 is read against the whole report.

Table 5: Issues of the UN report that this study does not cover.

Issue	UN §	Why it is outside this study
Privacy of the data	§16, §63	The study uses no personal data: every prompt is synthetic.
Intentionally discriminatory use: surveillance, disinformation	§19–20	The study tests advice given in good faith, not a deliberate use.
Remedies and appeals for victims	§23, §64, §68 g–h	A legal and institutional question; the study gives evidence, not a remedy.
Sector cases: facial recognition, predictive policing, recidivism, autonomous weapons, health scores, grading, AI detectors	§26–50	Other systems and other decisions; the study covers chat assistants giving safety advice.
Regulation: national, regional and international initiatives, the EU AI Act and its exceptions	§51–57	A review of law and policy, not of model behaviour.
Binding or voluntary standards, prohibitions, no exceptions	§68 a–c, e, i	Choices for legislators.
Consultation of the groups concerned, training of AI teams	§68 j, §69 b, e	Governance of how models are built.